Linux debugging, profiling, tracing and performance analysis training

On-site training, 3 days
Latest update: February 24, 2023

Title	**Linux debugging, profiling, tracing and performance analysis training**
Training objectives	• Be able to understand the main concepts of Linux that are relevant for performance analysis: process, threads, memory management, virtual memory, execution contexts, etc. • Be able to analyze why a system is loaded and what are the elements that contributes to this load using common Linux observability tools. • Be able to debug an userspace application using *gdb*, either live or after a crash, and analyze the contents of ELF binaries. • Be able to trace and profile a complete userspace application and its interactions with the Linux kernel in order to fix bugs using *strace, ltrace, perf* or *Callgrind*. • Be able to understand classical memory issues and analyze them using *valgrind, libefence* or *Massif*. • Be able to trace and profile the entire Linux system, using *perf, ftrace, kprobes, eBPF* tools, *kernelshark* or *LTTng* • Be able to debug Linux kernel issues: debug kernel crashes live or post-mortem, analyze memory issues at the kernel level, analyze locking issues, use kernel-level debuggers.
Duration	**Three** days - 24 hours (8 hours per day)
Pedagogics	• Lectures delivered by the trainer: 40% of the duration • Practical labs done by participants: 60% of the duration • Electronic copies of presentations, lab instructions and data files. They are freely available at https://bootlin.com/doc/training/debugging.
Trainer	Clément Léger https://bootlin.com/company/staff/clement-leger/
Language	Oral lectures: English, French. Materials: English.
Audience	Companies and engineers interested in debugging, profiling and tracing Linux systems and applications, to analyze and address performance or latency problems.

Prerequisites	• **Knowledge and practice of UNIX or GNU/Linux commands**: participants must be familiar with the Linux command line. Participants lacking experience on this topic should get trained by themselves, for example with our freely available on-line slides at bootlin.com/blog/command-line/. • **Minimal experience in embedded Linux development**: participants should have a minimal understanding of the architecture of embedded Linux systems: role of the Linux kernel vs. user-space, development of Linux user-space applications in C. Following Bootlin's *Embedded Linux* course at bootlin.com/training/embedded-linux/ allows to fulfill this pre-requisite. • **Minimal English language level: B1**, according to the *Common European Framework of References for Languages,* for our sessions in English. See bootlin.com/pub/training/cefr-grid.pdf for self-evaluation.
Required equipment	• Video projector • One PC computer on each desk (for one or two persons) with at least 8 GB of RAM, and Ubuntu Linux 22.04 installed in a **free partition of at least 30 GB** • Distributions others than Ubuntu Linux 22.04 are not supported, and using Linux in a virtual machine is not supported. • **Unfiltered and fast connection to Internet**: at least 50 Mbit/s of download bandwidth, and no filtering of web sites or protocols. • **PC computers with valuable data must be backed up** before being used in our sessions.
Certificate	Only the participants who have attended all training sessions, and who have scored over 50% of correct answers at the final evaluation will receive a training certificate from Bootlin.
Disabilities	Participants with disabilities who have special needs are invited to contact us at *training@bootlin.com* to discuss adaptations to the training course.

Hardware in practical labs

The hardware platform used for the practical labs of this training session is the **STMicroelectronics STM32MP157D-DK1 Discovery board** board, which features:

- STM32MP157D (dual Cortex-A7) CPU from STMicroelectronics
- USB powered
- 512 MB DDR3L RAM
- Gigabit Ethernet port
- 4 USB 2.0 host ports
- 1 USB-C OTG port
- 1 Micro SD slot
- On-board ST-LINK/V2-1 debugger
- Arduino Uno v3-compatible headers
- Audio codec
- Misc: buttons, LEDs

Day 1 - Morning

Lecture - Linux application stack

- Global picture: understanding the general architecture of a Linux system, overview of the major components.
- What is the difference between a process and a thread, how applications run concurrently.
- ELF files and associated analysis tools.
- Userspace application memory layout (heap, stack, shared libraries mappings, etc).
- MMU and memory management: physical/virtual address spaces.
- Kernel context switching and scheduling
- Kernel execution contexts: kernel threads, workqueues, interrupt, threaded interrupts, softirq

Lecture - Common analysis & observability tools

- Analyzing an ELF file with GNU binary utilities (*objdump, addr2line*).
- Tools to monitor a Linux system: processes, memory usage and mapping, resources.
- Using *vmstat, iostat, ps, top, iotop, free* and understanding the metrics they provide.
- Pseudo filesystems: *procfs, sysfs* and *debugfs*.

Day 1 - Afternoon

**Lab - Check what is running on a system and its load**

- Observe running processes using *ps* and *top*.
- Check memory allocation and mapping with *procfs* and *pmap*.
- Monitor other resources usage using *iostat, vmstat* and *netstat*.

Lecture - Debugging an application

- Using *gdb* on a live process.
- Understanding compiler optimizations impact on debuggability.
- Postmortem diagnostic using core files.
- Remote debugging with *gdbserver*.
- Extending *gdb* capabilities using python scripting

Lab - Solving an application crash

- Analysis of compiled C code with compiler-explorer to understand optimizations.
- Managing *gdb* from the command line, then from an IDE.
- Using *gdb* Python scripting capabilities.
- Debugging a crashed application using a core-dump with *gdb*.
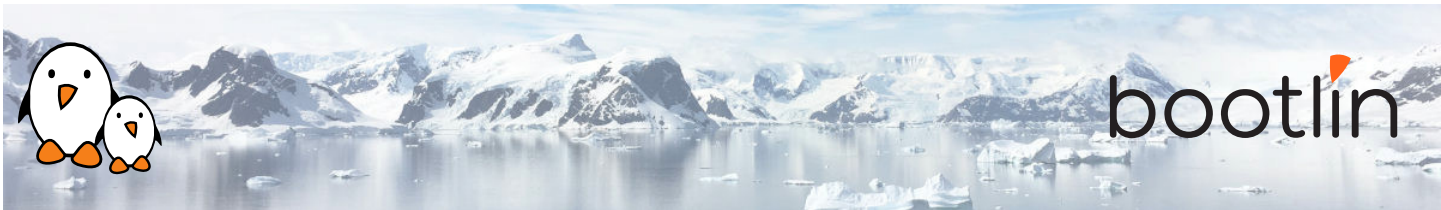
Day 2 - Morning

Lecture - Tracing an application

- Tracing system calls with *strace*.
- Tracing library calls with *ltrace*.
- Overloading library functions using *LD_-PRELOAD*.

Lab – Debugging application issues

- Analyze dynamic library calls from an application using *ltrace*.
- Overloading library functions using *LD_-PRELOAD*.
- Analyzing an application system calls using *strace*.

Lecture - Memory issues

- Usual memory issues: buffer overflow, segmentation fault, memory leaks, heap-stack collision.
- Memory corruption tooling, *valgrind, libefence*, etc.
- Heap profiling using *Massif* and *heaptrack*

Lab – Debugging memory issues

- Memory leak and misbehavior detection with *valgrind* and *vgdb*.
- Visualizing application heap using *Massif*.

Day 2 - Afternoon

Lecture – Application profiling

- Performances issues.
- Gathering profiling data with *perf*.
- Analyzing an application callgraph using *Callgrind* and *KCachegrind*.
- Interpreting the data recorded by *perf*.

Lab - Application profiling

- Profiling an application with *Callgrind/KCachegrind*.
- Analyzing application performance with *perf*.
- Generating a flamegraph using *FlameGraph*.
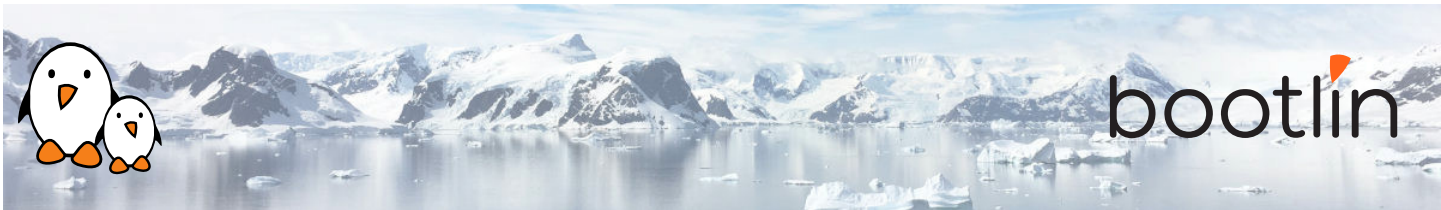
Day 3 - Morning

Lecture - System wide profiling and tracing	Lab - System wide profiling and tracing
• System wide profiling using *perf*. • Using *kprobes* to hook on kernel code without recompiling. • *eBPF* tools (*bcctools, bpftrace,* etc) for complex tracing scenarios. • Application and kernel tracing and visualization using *ftrace, kernelshark* or *LTTng*	• System profiling with *perf*. • IRQ latencies using *ftrace*. • Tracing and visualizing system activity using *kernelshark* or *LTTng*

Day 3 - Afternoon

Lecture - Kernel debugging	Lab - Kernel debugging
• Kernel compilation results (`vmlinux`, `System.map`). • Understanding and configuring kernel *oops* behavior. • Post mortem analysis using kernel crash dump with *crash*. • Memory issues (*KASAN, UBSAN, Kmemleak*). • Debugging the kernel using *KGDB* and *KDB*. • Kernel locking debug configuration options (lockdep). • Other kernel configuration options that are useful for debug.	• Analyzing an *oops* after using a faulty module with *obdjump* and *addr2line*. • Debugging a deadlock problem using *PROVE_-LOCKING* options. • Detecting undefined behavior with *UBSAN* in kernel code. • Find a module memory leak using *kmemleak*. • Debugging a module with *KGDB*.